
Transductive Zero-Shot & Few-Shot Adaptation of Large Medical Vision Language Models



Ismail Ben Ayed



Julio Silva-Rodriguez



Maxime Zanella



Yunshi Huang



Jose Dolz



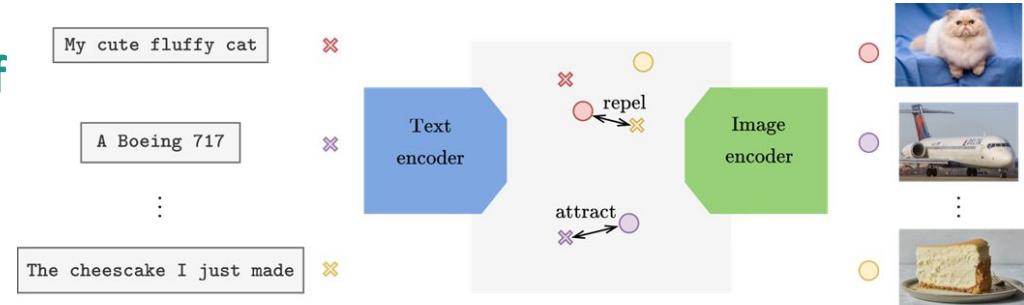
Houda Bahig

Fereshteh Shakeri
Jan 2026

Lots of different Vision-Language Models out there

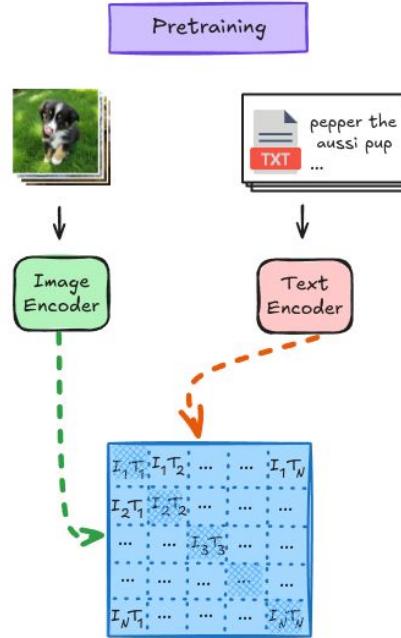
Among them...

- **CLIP: generic text-image pairs of internet**



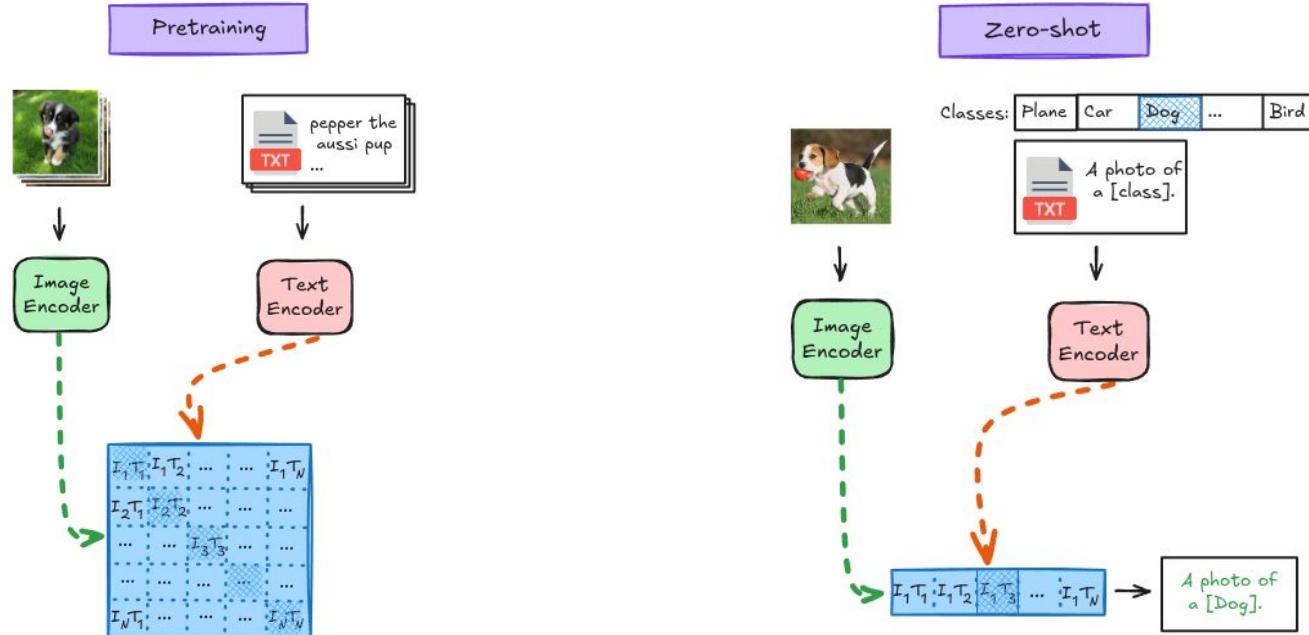
CLIP: Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021, July). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748-8763). PMLR.

Generalist Vision-language Models (CLIP)



Radford, Alec, et al. "Learning transferable visual models from natural language supervision." PMLR, 2021.

Generalist Vision-language Models (CLIP)

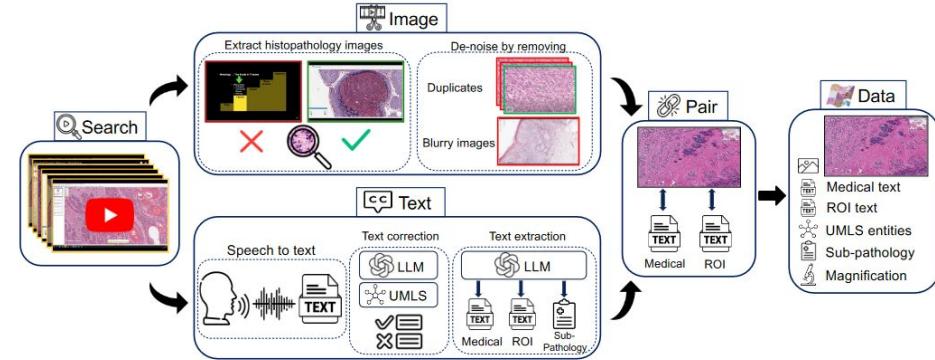


Radford, Alec, et al. "Learning transferable visual models from natural language supervision." PMLR, 2021.

Lots of different Vision-Language Models out there

Among them...

- CLIP: generic text-image pairs of internet
- Quilt: from Youtube



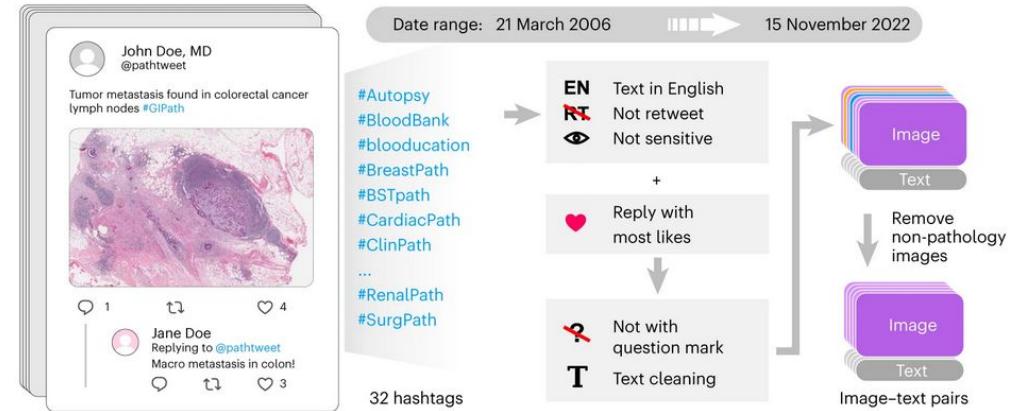
CLIP: Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021, July). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748-8763). PMLR.

Quilt: Ikezogwo, W., Seyfioglu, S., Ghezloo, F., Geva, D., Sheikh Mohammed, F., Anand, P. K., ... & Shapiro, L. (2024). Quilt-1m: One million image-text pairs for histopathology. *Advances in neural information processing systems*, 36.

Lots of different Vision-Language Models out there

Among them...

- CLIP: generic text-image pairs of internet
- Quilt: from Youtube
- PLIP: from Twitter



CLIP: Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021, July). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748-8763). PMLR.

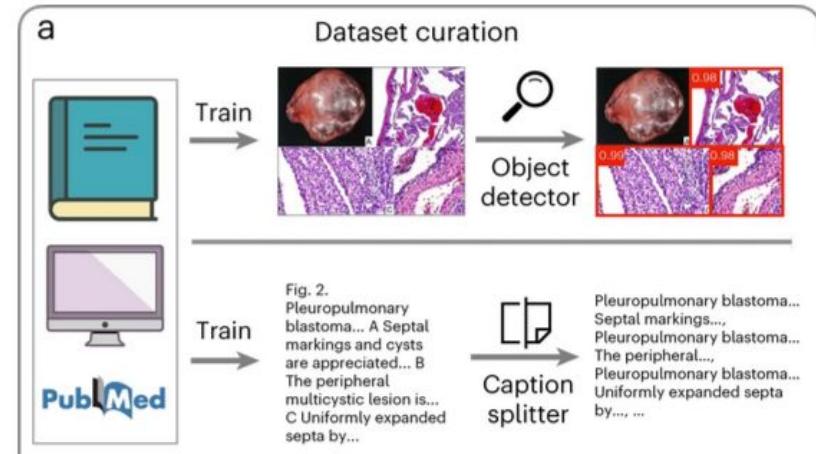
Quilt: Ikezogwo, W., Seyfoglou, S., Ghezlo, F., Geva, D., Sheikh Mohammed, F., Anand, P. K., ... & Shapiro, L. (2024). Quilt-1m: One million image-text pairs for histopathology. *Advances in neural information processing systems*, 36.

PLIP: Huang, Z., Bianchi, F., Yuksekgonul, M., Montine, T. J., & Zou, J. (2023). A visual-language foundation model for pathology image analysis using medical twitter. *Nature medicine*, 29(9), 2307-2316.

Lots of different Vision-Language Models out there

Among them...

- CLIP: generic text-image pairs of internet
- Quilt: from Youtube
- PLIP: from Twitter
- CONCH : from PubMed



CLIP: Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021, July). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748-8763). PMLR.

Quilt: Ikezogwo, W., Seyfoglou, S., Ghezlo, F., Geva, D., Sheikh Mohammed, F., Anand, P. K., ... & Shapiro, L. (2024). Quilt-1m: One million image-text pairs for histopathology. *Advances in neural information processing systems*, 36.

PLIP: Huang, Z., Bianchi, F., Yuksekgonul, M., Montine, T. J., & Zou, J. (2023). A visual-language foundation model for pathology image analysis using medical twitter. *Nature medicine*, 29(9), 2307-2316.

CONCH: Lu, M. Y., Chen, B., Williamson, D. F., Chen, R. J., Liang, I., Ding, T., ... & Mahmood, F. (2024). A visual-language foundation model for computational pathology. *Nature Medicine*, 30(3), 863-874.

Lots of different Vision-Language Models out there

Among them...

- CLIP: generic text-image pairs of internet
- Quilt: from Youtube
- PLIP: from Twitter
- CONCH : from PubMed

How to leverage their representational power?

...

Boosting Vision-Language Models for Histopathology Classification:
Predict all at once

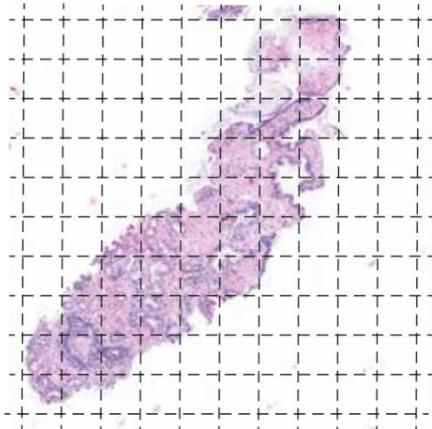
CLIP: Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021, July). Learning transferable visual models from natural language supervision. In *International conference on machine learning* (pp. 8748-8763). PMLR.

Quilt: Ikezogwo, W., Seyfoglou, S., Ghezlo, F., Geva, D., Sheikh Mohammed, F., Anand, P. K., ... & Shapiro, L. (2024). Quilt-1m: One million image-text pairs for histopathology. *Advances in neural information processing systems*, 36.

PLIP: Huang, Z., Bianchi, F., Yuksekgonul, M., Montine, T. J., & Zou, J. (2023). A visual–language foundation model for pathology image analysis using medical twitter. *Nature medicine*, 29(9), 2307-2316.

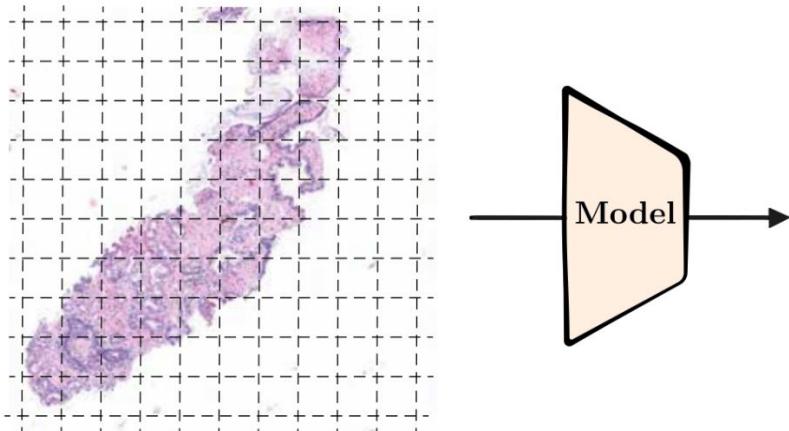
CONCH: Lu, M. Y., Chen, B., Williamson, D. F., Chen, R. J., Liang, I., Ding, T., ... & Mahmood, F. (2024). A visual-language foundation model for computational pathology. *Nature Medicine*, 30(3), 863-874.

Inductive inference



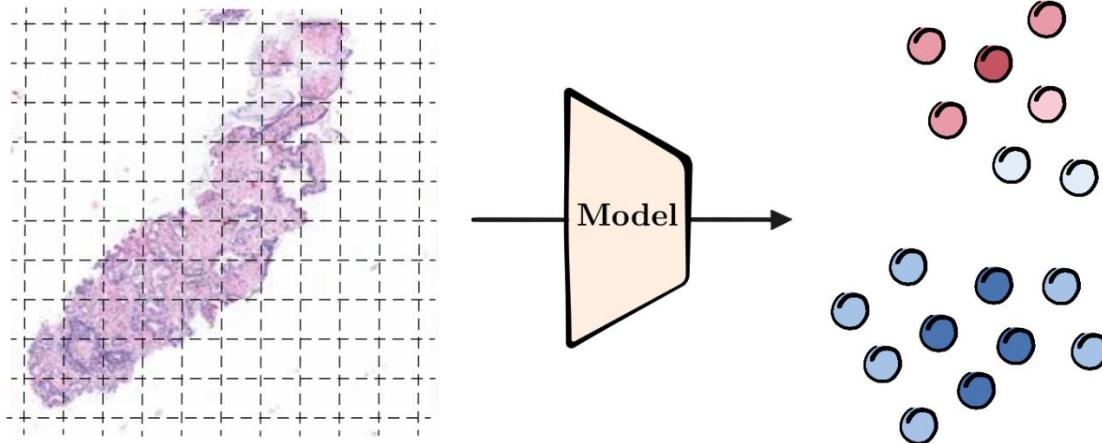
Large whole slides divided per patches, large-scale database, ...

Inductive inference



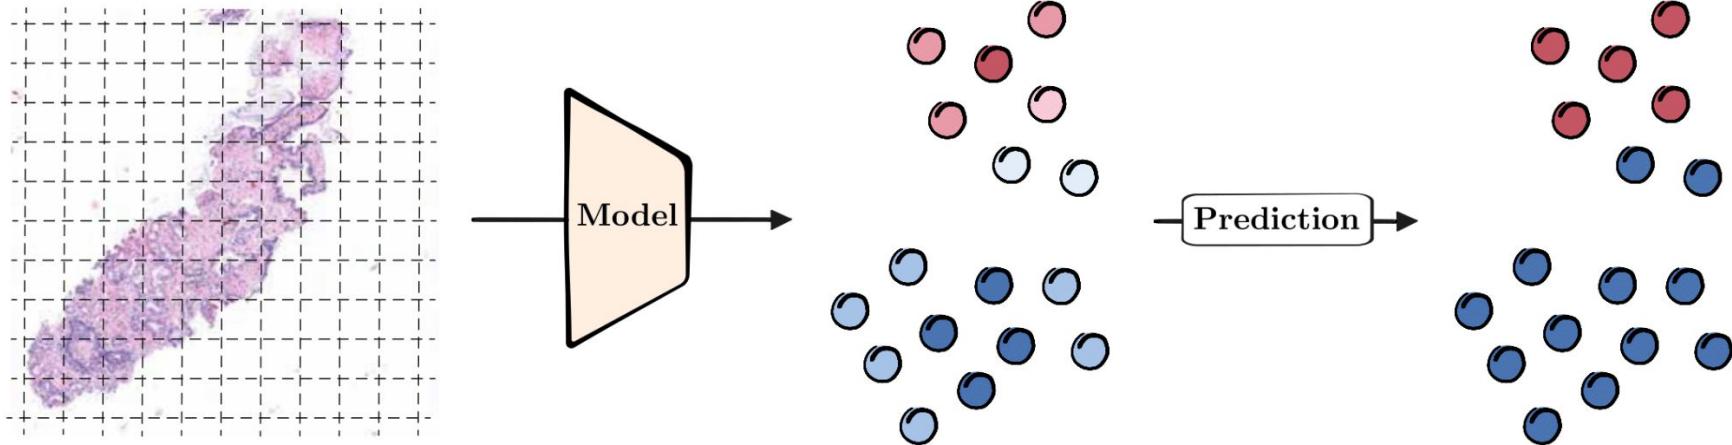
(1) If we have **enough annotated data**, we can train a Model and use it at inference time...

Inductive inference



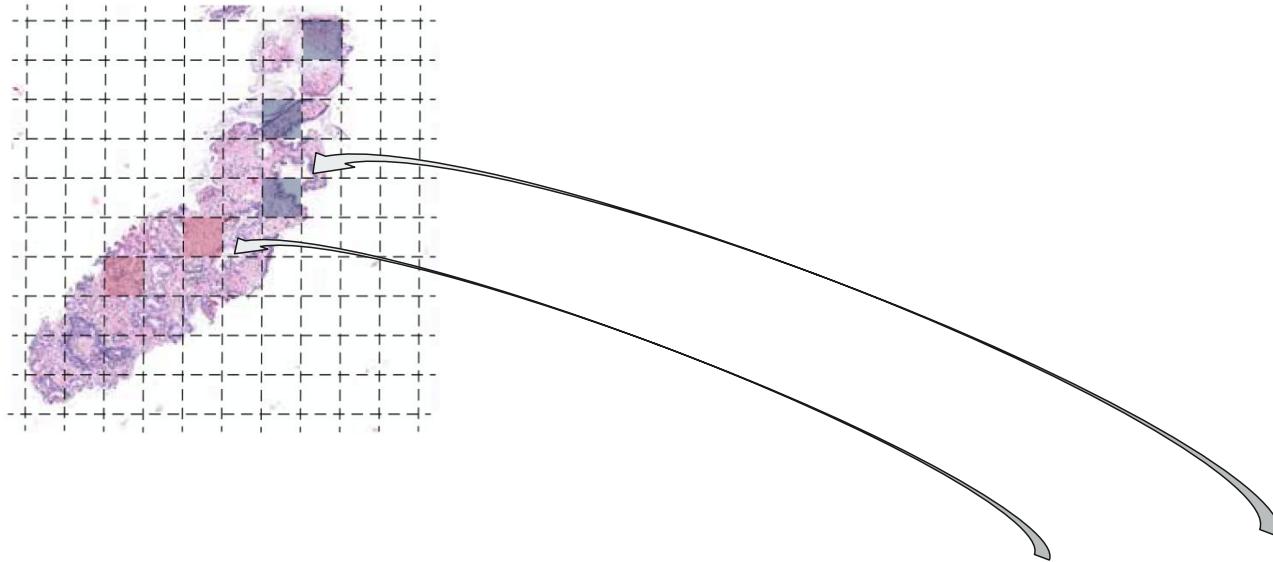
... which gives us more or less confident predictions ...

Inductive inference



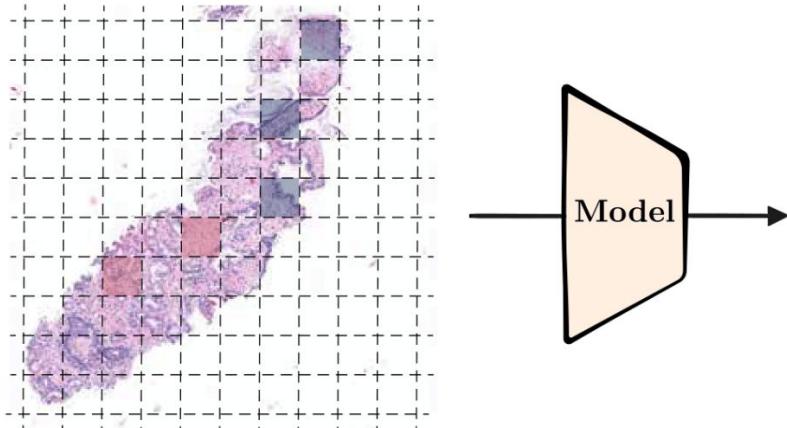
... and we take the most probable class for **each patch independently**.

Few-shot transductive inference



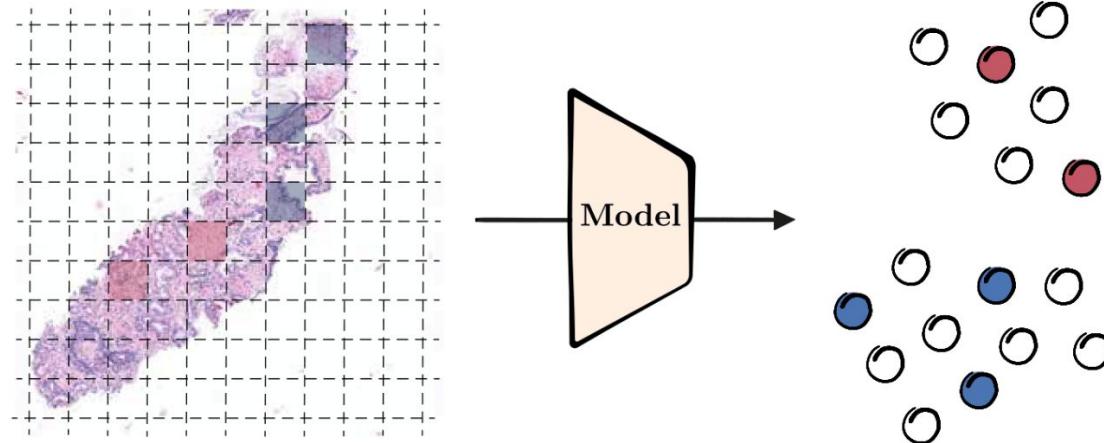
Large whole slides divided per patches, large-scale database, ... + a few annotated patches

Few-shot transductive inference



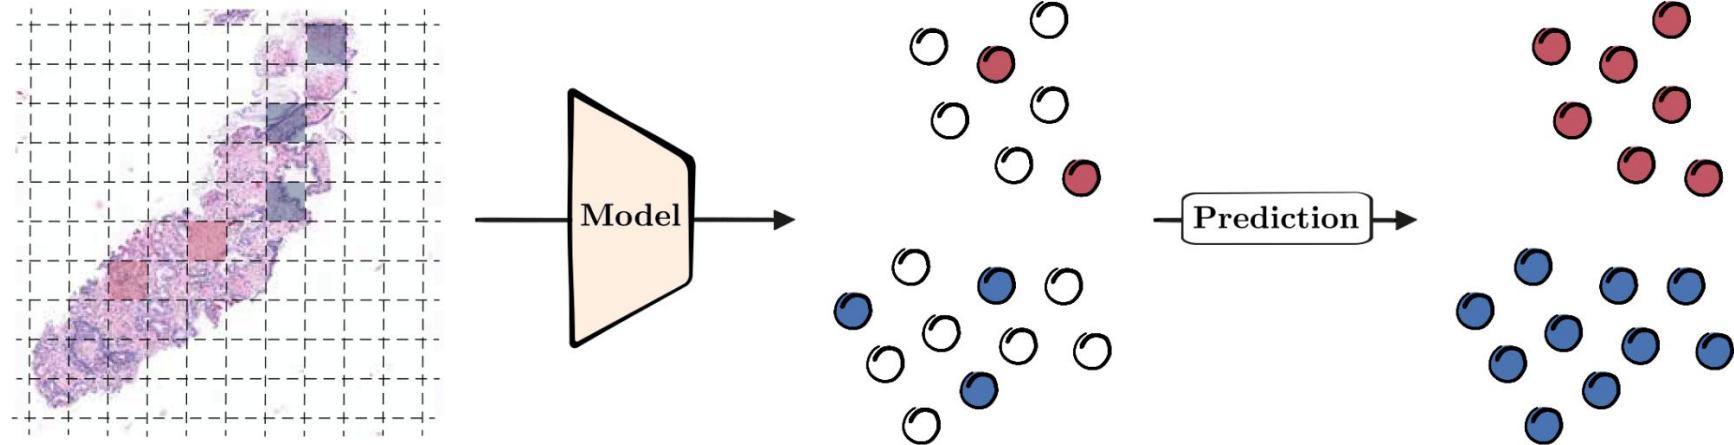
(2) We can use a good embedder (even self-supervised!) ...

Few-shot transductive inference



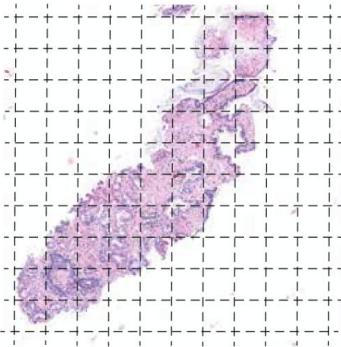
... which gives a good representation ...

Few-shot transduction inference



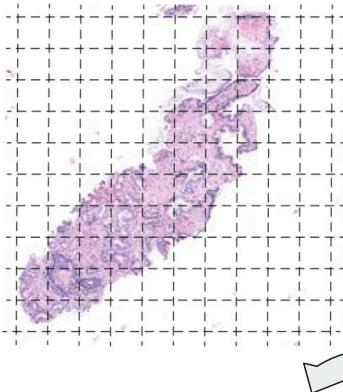
... we can leverage **data structure** and the few-labeled patches !

Unsupervised transductive inference

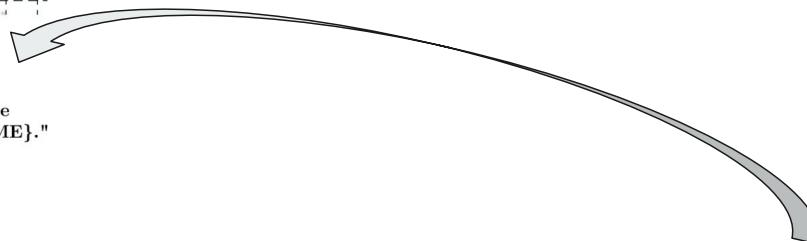


Large whole slides divided per patches, large-scale database, ... +~~a few annotated patches~~

Unsupervised transductive inference

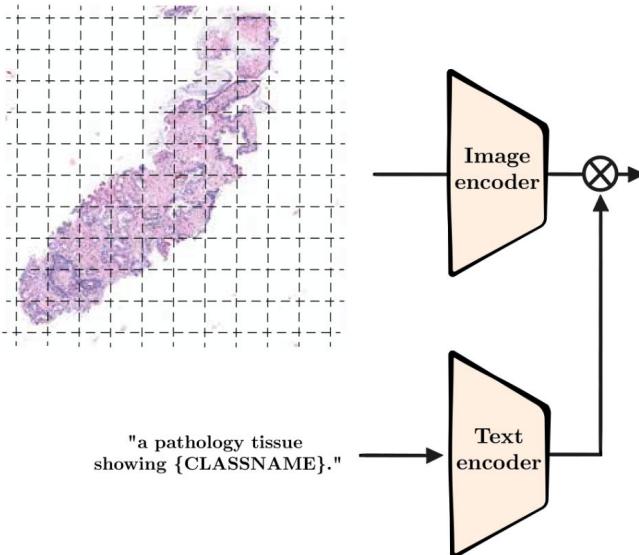


"a pathology tissue
showing {CLASSNAME}."



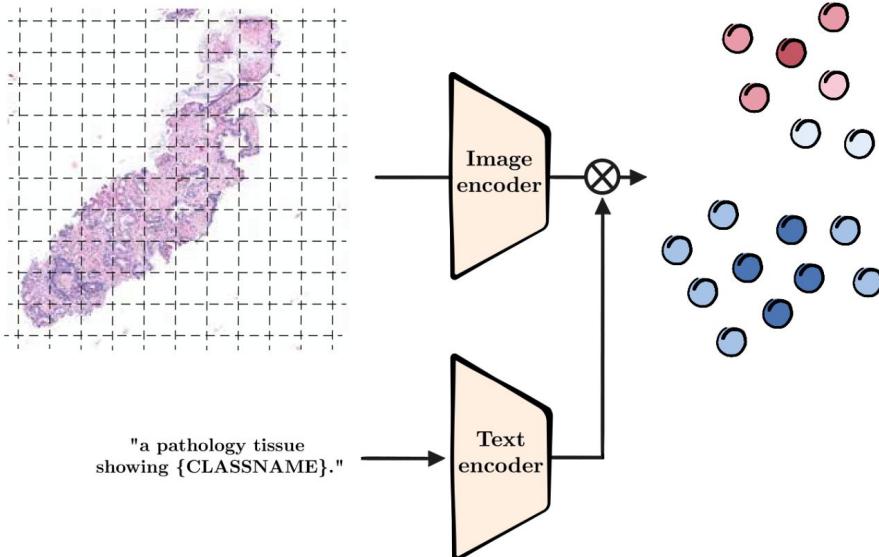
Large whole slides divided per patches, large-scale database, ... + **text description for each class**

Unsupervised transductive inference



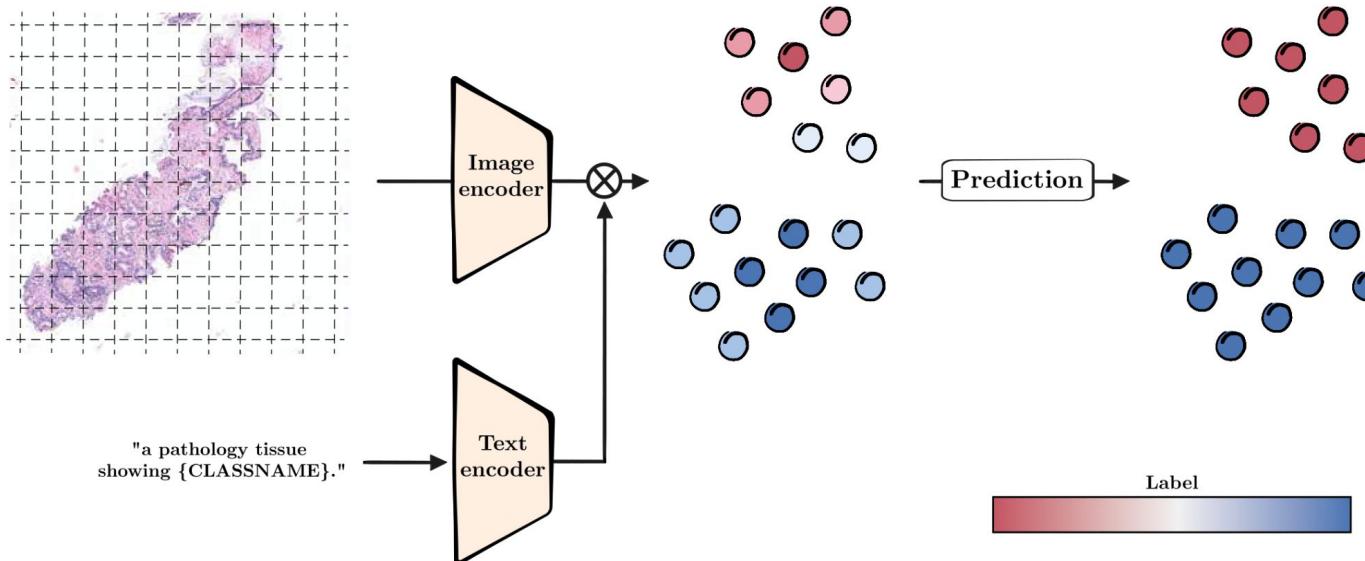
(3) we leverage **Vision-Language Models (VLMs)** trained on large-scale unsupervised text-image data...

Unsupervised transductive inference



... text-image similarities give us (noisy) **zero-shot predictions** for each patch ...

Unsupervised transductive inference



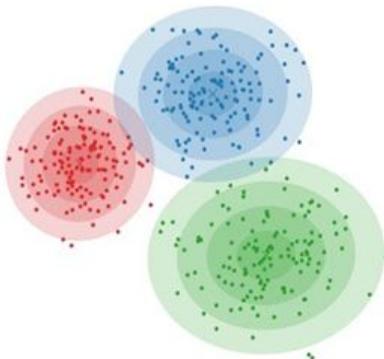
... these **zero-shot predictions** can be used in combination with the **data structure!**

Our algorithm (in short)

$$\mathcal{L}_{\text{Zero-Shot}}(\mathbf{z}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \underbrace{-\frac{1}{|\mathcal{Q}|} \sum_{i \in \mathcal{Q}} \mathbf{z}_i^\top \log(\mathbf{p}_i)}_{\text{GMM clustering}} - \underbrace{\sum_{i \in \mathcal{D}} \sum_{j \in \mathcal{D}} w_{ij} \mathbf{z}_i \mathbf{z}_j}_{\text{Laplacian reg.}} + \underbrace{\sum_{i \in \mathcal{Q}} \text{KL}_\lambda(\mathbf{z}_i || \mathbf{j})}_{\text{Text knowledge}}$$

Our algorithm (in short)

$$\mathcal{L}_{\text{Zero-Shot}}(\mathbf{z}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \underbrace{-\frac{1}{|\mathcal{Q}|} \sum_{i \in \mathcal{Q}} \mathbf{z}_i^\top \log(\mathbf{p}_i)}_{\text{GMM clustering}} - \underbrace{\sum_{i \in \mathcal{D}} \sum_{j \in \mathcal{D}} w_{ij} \mathbf{z}_i \mathbf{z}_j}_{\text{Laplacian reg.}} + \underbrace{\sum_{i \in \mathcal{Q}} \text{KL}_\lambda(\mathbf{z}_i || \hat{\mathbf{y}}_i)}_{\text{Text knowledge}}$$



We model features as a balanced Gaussian Mixture Model (**GMM**).

Our algorithm (in short)

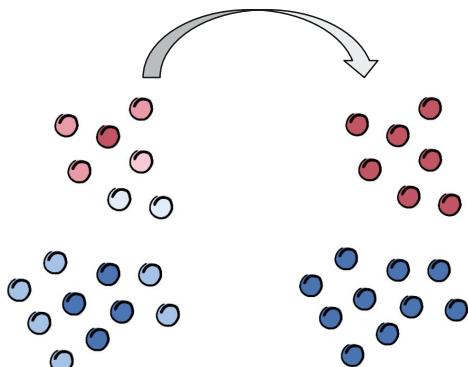
$$\mathcal{L}_{\text{Zero-Shot}}(\mathbf{z}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \underbrace{-\frac{1}{|\mathcal{Q}|} \sum_{i \in \mathcal{Q}} \mathbf{z}_i^\top \log(\mathbf{p}_i)}_{\text{GMM clustering}} - \underbrace{\sum_{i \in \mathcal{D}} \sum_{j \in \mathcal{D}} w_{ij} \mathbf{z}_i \mathbf{z}_j}_{\text{Laplacian reg.}} + \underbrace{\sum_{i \in \mathcal{Q}} \text{KL}_\lambda(\mathbf{z}_i || \hat{\mathbf{y}}_i)}_{\text{Text knowledge}}$$



Similar patches should have similar predictions!

Our algorithm (in short)

$$\mathcal{L}_{\text{Zero-Shot}}(\mathbf{z}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \underbrace{-\frac{1}{|\mathcal{Q}|} \sum_{i \in \mathcal{Q}} \mathbf{z}_i^\top \log(\mathbf{p}_i)}_{\text{GMM clustering}} - \underbrace{\sum_{i \in \mathcal{D}} \sum_{j \in \mathcal{D}} w_{ij} \mathbf{z}_i \mathbf{z}_j}_{\text{Laplacian reg.}} + \underbrace{\sum_{i \in \mathcal{Q}} \text{KL}_\lambda(\mathbf{z}_i || \hat{\mathbf{y}}_i)}_{\text{Text knowledge}}$$



New predictions should remain **as close as possible to the initial (zero-shot) ones** (while minimizing the 2 other terms!)

Our algorithm (in short)

Our algorithm (in short)

Solving procedure of 3 update equations
 → just a **few lines of code!**

```

1 function Histo-TransCLIP( $f$ ,  $t$ ,  $\tau$ )
2   // Text-based pseudo-labels  $\hat{y}$ 
3    $\hat{y}_i = \text{softmax}(\tau f_i^T t)$     $\forall i$ 
4   // Initialize  $z$ ,  $\mu$ ,  $\Sigma$ 
5    $z_i = \hat{y}_i$     $\forall i$ 
6    $\mu_k = \text{top\_confident\_average}(f, \hat{y})$     $\forall k$ 
7    $\text{diag}(\Sigma) = \frac{1}{n\_features}$ 
8   // Iterative procedure
9   while not_converged do
10    for  $l = 1:\dots$  do
11       $z_i^{(l+1)} = \frac{\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)})}{(\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)}))^T \mathbb{1}_K}$     $\forall i$ 
12       $\mu_k = \frac{\sum_{i \in Q} z_{i,k} f_i}{\sum_{i \in Q} z_{i,k}}$     $\forall k$ 
13       $\text{diag}(\Sigma) = \frac{1}{|Q|} \sum_{i \in Q} \sum_k z_{i,k} (f_i - \mu_k)^2$ 
14
15 return  $z$ 

```

Our algorithm (in short)

Solving procedure of 3 update equations
 → just a **few lines of code!**

Really fast to solve
 → just a **few seconds for 100,000 patches**

```

1 function Histo-TransCLIP( $f$ ,  $t$ ,  $\tau$ )
2   // Text-based pseudo-labels  $\hat{y}$ 
3    $\hat{y}_i = \text{softmax}(\tau f_i^T t)$     $\forall i$ 
4   // Initialize  $z$ ,  $\mu$ ,  $\Sigma$ 
5    $z_i = \hat{y}_i$     $\forall i$ 
6    $\mu_k = \text{top\_confident\_average}(f, \hat{y})$     $\forall k$ 
7    $\text{diag}(\Sigma) = \frac{1}{n\_features}$ 
8   // Iterative procedure
9   while not_converged do
10    for  $l = 1:\dots$  do
11       $z_i^{(l+1)} = \frac{\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)})}{(\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)}))^T \mathbb{1}_K}$     $\forall i$ 
12       $\mu_k = \frac{\sum_{i \in Q} z_{i,k} f_i}{\sum_{i \in Q} z_{i,k}}$     $\forall k$ 
13       $\text{diag}(\Sigma) = \frac{1}{|Q|} \sum_{i \in Q} \sum_k z_{i,k} (f_i - \mu_k)^2$ 
14
15 return  $z$ 

```

Our algorithm (in short)

Solving procedure of 3 update equations
 → just a **few lines of code!**

Really fast to solve
 → just a **few seconds for 100,000 patches**

#Patches	Features	Histo-TransCLIP
10^2	~ 1 sec.	~ 0.1 sec.
10^3	~ 4 sec.	~ 0.2 sec.
10^4	~ 28 sec.	~ 0.4 sec.
10^5	~ 5 min.	~ 6 sec.

```

1 function Histo-TransCLIP( $f$ ,  $t$ ,  $\tau$ )
2   // Text-based pseudo-labels  $\hat{y}$ 
3    $\hat{y}_i = \text{softmax}(\tau f_i^T t)$     $\forall i$ 
4   // Initialize  $z$ ,  $\mu$ ,  $\Sigma$ 
5    $z_i = \hat{y}_i$     $\forall i$ 
6    $\mu_k = \text{top\_confident\_average}(f, \hat{y})$     $\forall k$ 
7    $\text{diag}(\Sigma) = \frac{1}{n\_features}$ 
8   // Iterative procedure
9   while not_converged do
10    for  $l = 1:\dots$  do
11       $z_i^{(l+1)} = \frac{\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)})}{(\hat{y}_i \odot \exp(\log(p_i) + \sum_{j \in Q} w_{ij} z_j^{(l)}))^T \mathbb{1}_K}$     $\forall i$ 
12       $\mu_k = \frac{\sum_{i \in Q} z_{i,k} f_i}{\sum_{i \in Q} z_{i,k}}$     $\forall k$ 
13       $\text{diag}(\Sigma) = \frac{1}{|Q|} \sum_{i \in Q} \sum_k z_{i,k} (f_i - \mu_k)^2$ 
14
15   return  $z$ 

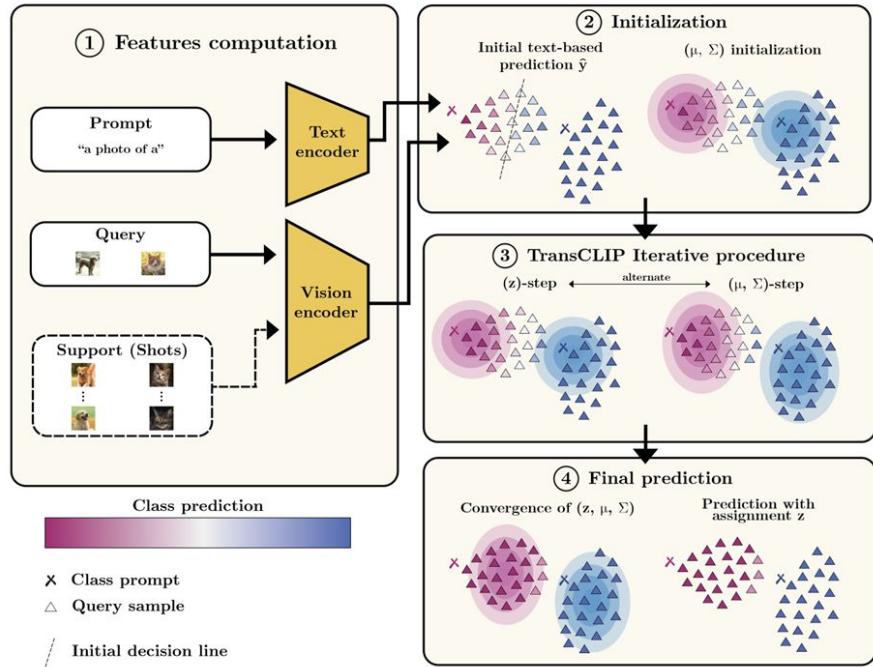
```

Our algorithm (in short)

Solving procedure of 3 update equations
 → just a **few lines of code!**

Really fast to solve
 → just a **few seconds for 100,000 patches!**

#Patches	Features	Histo-TransCLIP
10^2	~ 1 sec.	~ 0.1 sec.
10^3	~ 4 sec.	~ 0.2 sec.
10^4	~ 28 sec.	~ 0.4 sec.
10^5	~ 5 min.	~ 6 sec.

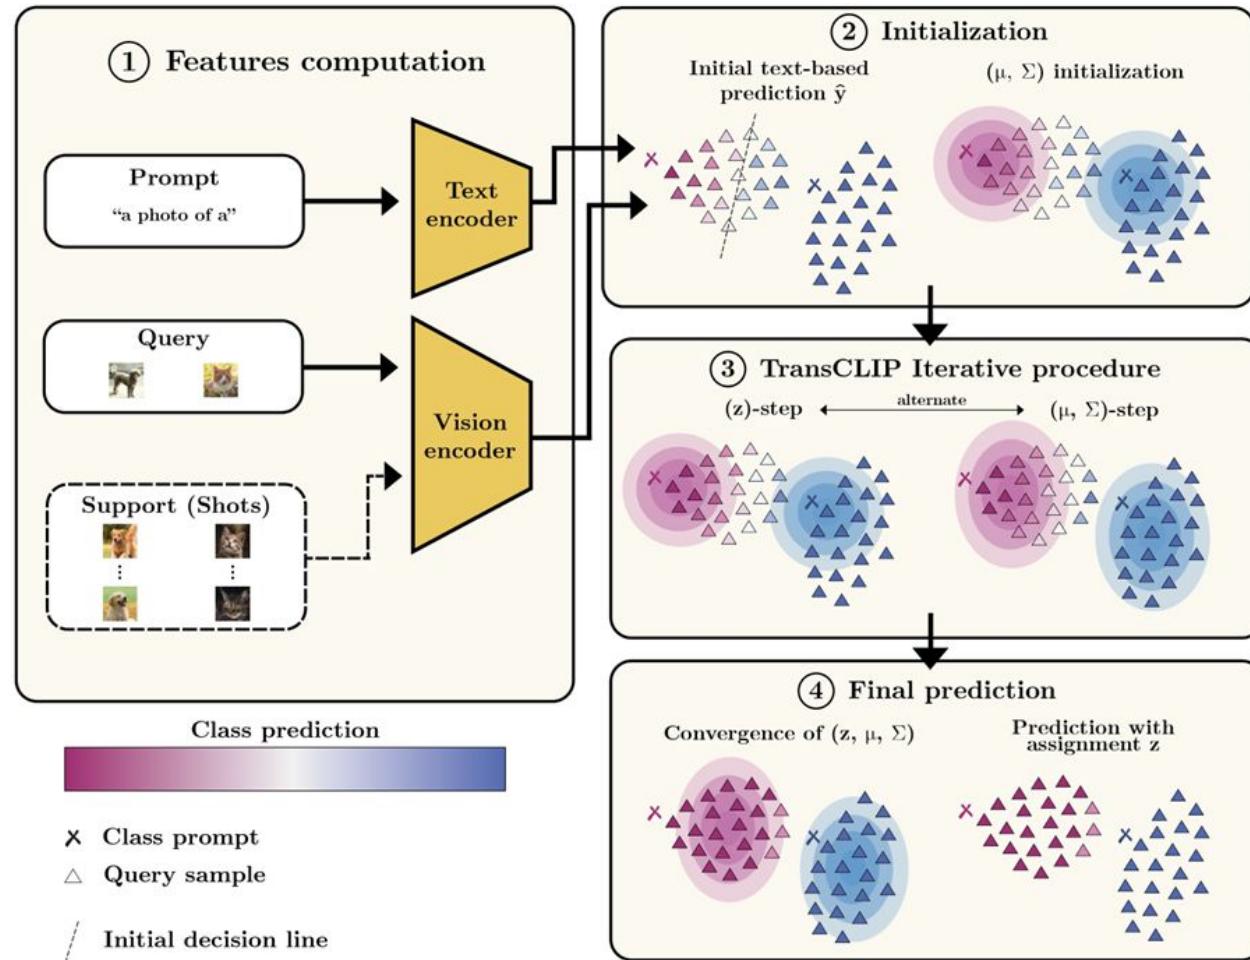


Our algorithm

Solving problems
→ just a few

Really fast time
→ just a few

#I



Results

Dataset	Method	Model				
		CLIP	Quilt-B16	Quilt-B32	PLIP	CONCH
<i>SICAP-MIL</i>	Zero-shot	29.85	40.44	35.04	46.84	27.71
	Histo-TransCLIP	24.72	58.49	28.18	53.23	32.58
<i>LC(Lung)</i>	Zero-shot	31.46	43.00	76.24	84.96	84.81
	Histo-TransCLIP	25.62	50.53	93.93	93.80	96.29
<i>SKINCANCER</i>	Zero-shot	4.20	15.38	39.71	22.90	58.53
	Histo-TransCLIP	11.46	33.33	48.80	36.72	66.22
<i>NCT-CRC</i>	Zero-shot	25.39	29.61	53.73	63.17	66.27
	Histo-TransCLIP	39.61	48.40	58.13	77.53	70.36
<i>Average</i>	Zero-shot	22.73	32.1	51.18	54.47	59.33
	Histo-TransCLIP	25.35	47.69	57.26	65.32	66.36
	$\Delta_{\text{transductive}}$	+2.62	+15.59	+6.08	+10.85	+7.03

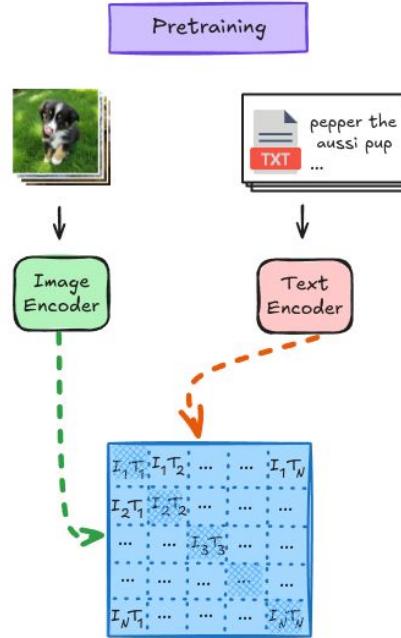
Results: generic pre-training

Dataset	Method	Model				
		CLIP	Quilt-B16	Quilt-B32	PLIP	CONCH
<i>SICAP-MIL</i>	Zero-shot	29.85	40.44	35.04	46.84	27.71
	Histo-TransCLIP	24.72	58.49	28.18	53.23	32.58
<i>LC(Lung)</i>	Zero-shot	31.46	43.00	76.24	84.96	84.81
	Histo-TransCLIP	25.62	50.53	93.93	93.80	96.29
<i>SKINCANCER</i>	Zero-shot	4.20	15.38	39.71	22.90	58.53
	Histo-TransCLIP	11.46	33.33	48.80	36.72	66.22
<i>NCT-CRC</i>	Zero-shot	25.39	29.61	53.73	63.17	66.27
	Histo-TransCLIP	39.61	48.40	58.13	77.53	70.36
<i>Average</i>	Zero-shot	22.73	32.1	51.18	54.47	59.33
	Histo-TransCLIP	25.35	47.69	57.26	65.32	66.36
	$\Delta_{\text{transductive}}$	+2.62	+15.59	+6.08	+10.85	+7.03

Results: histopathology pre-training

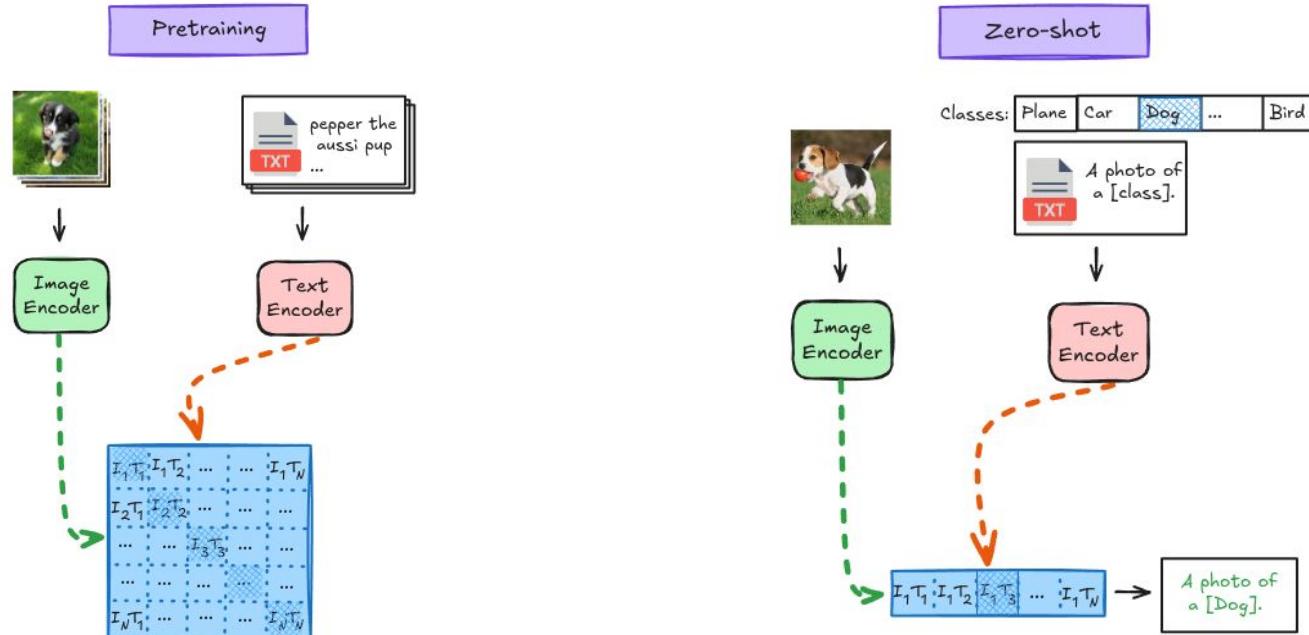
Dataset	Method	Model				
		CLIP	Quilt-B16	Quilt-B32	PLIP	CONCH
<i>SICAP-MIL</i>	Zero-shot	29.85	40.44	35.04	46.84	27.71
	Histo-TransCLIP	24.72	58.49	28.18	53.23	32.58
<i>LC(Lung)</i>	Zero-shot	31.40	43.00	76.24	84.96	84.81
	Histo-TransCLIP	25.63	50.53	93.93	93.80	96.29
<i>SKINCANCER</i>	Zero-shot	4.20	15.38	39.71	22.90	58.53
	Histo-TransCLIP	11.43	33.33	48.80	36.72	66.22
<i>NCT-CRC</i>	Zero-shot	25.39	29.61	53.73	63.17	66.27
	Histo-TransCLIP	39.61	48.40	58.13	77.53	70.36
<i>Average</i>	Zero-shot	22.73	32.1	51.18	54.47	59.33
	Histo-TransCLIP	25.35	47.69	57.26	65.32	66.36
	$\Delta_{\text{transductive}}$	+2.62	+15.59	+6.08	+10.85	+7.03

Generalist Vision-language Models (CLIP)



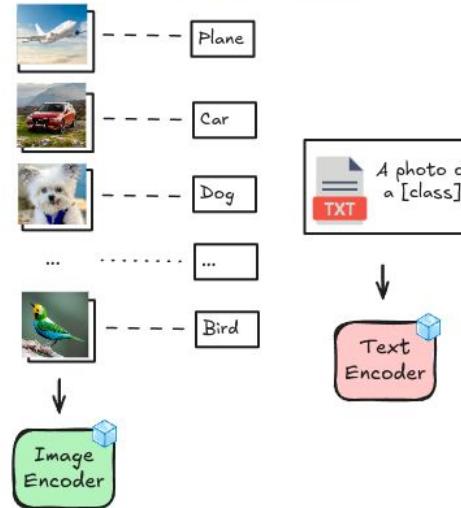
Radford, Alec, et al. "Learning transferable visual models from natural language supervision." PMLR, 2021.

Generalist Vision-language Models (CLIP)



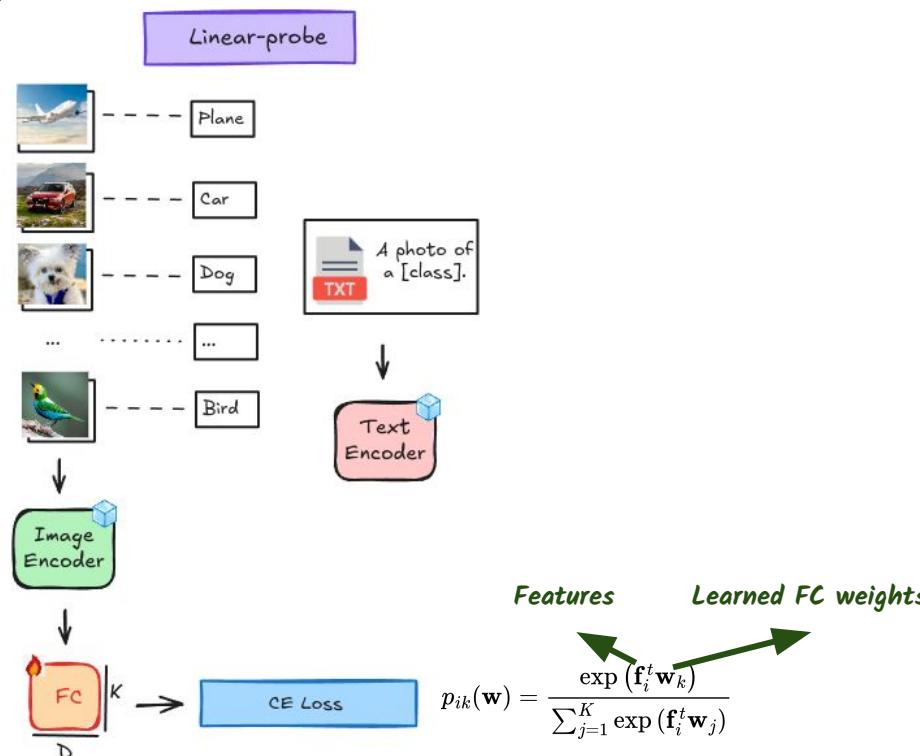
Radford, Alec, et al. "Learning transferable visual models from natural language supervision." PMLR, 2021.

Few-Shot Adaptation (linear probe)

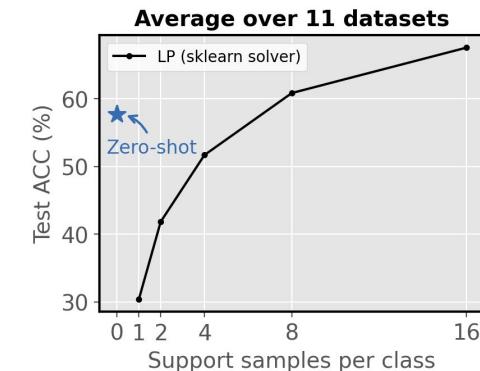
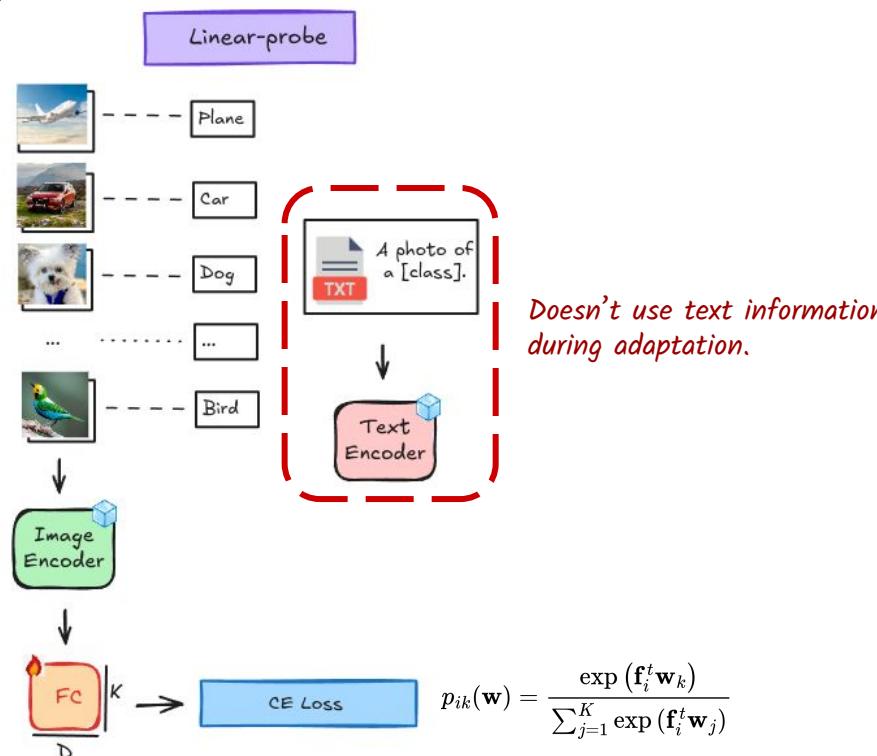


Radford, Alec, et al. "Learning transferable visual models from natural language supervision." PMLR, 2021.

Few-Shot Adaptation (linear probe)

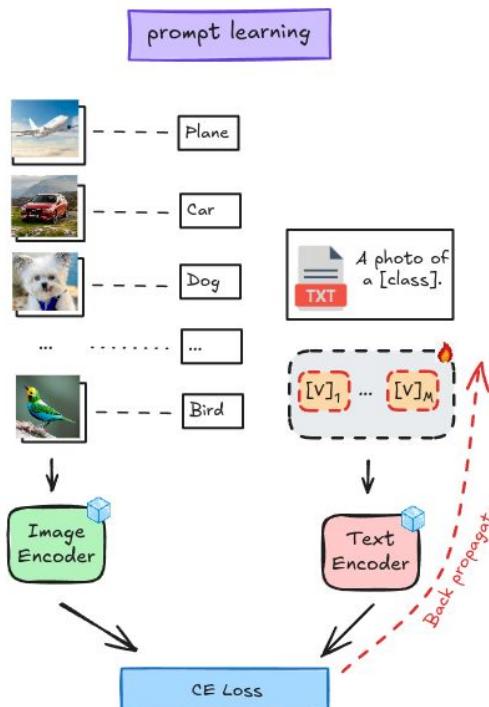


Few-Shot Adaptation (linear probe)



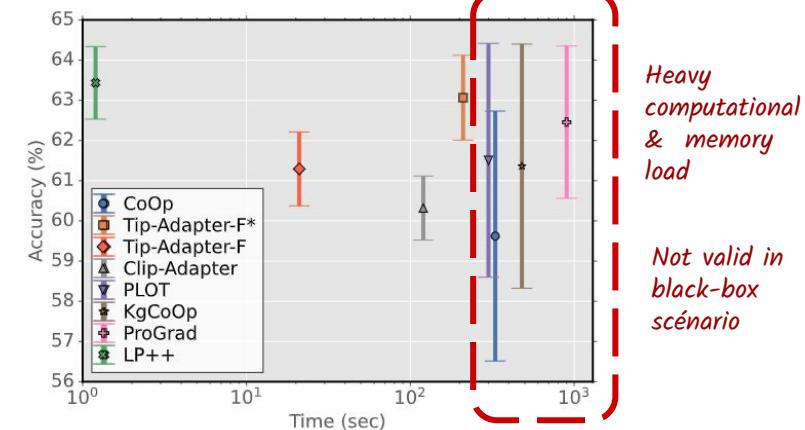
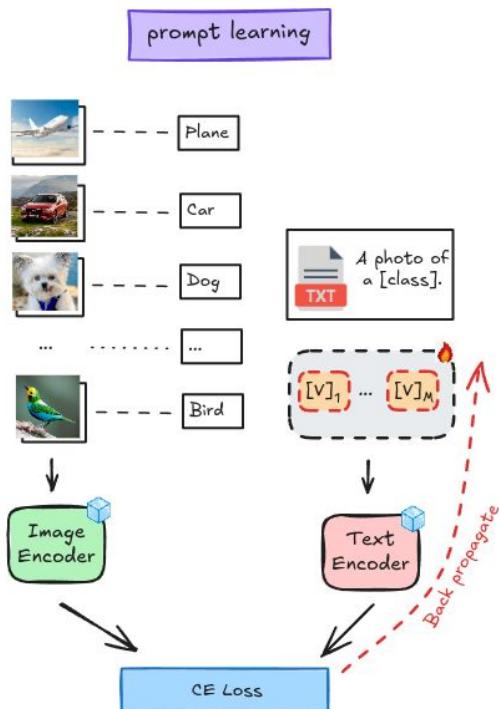
Linear Probe below Zero shot!!

Few-Shot Adaptation (prompt learning)



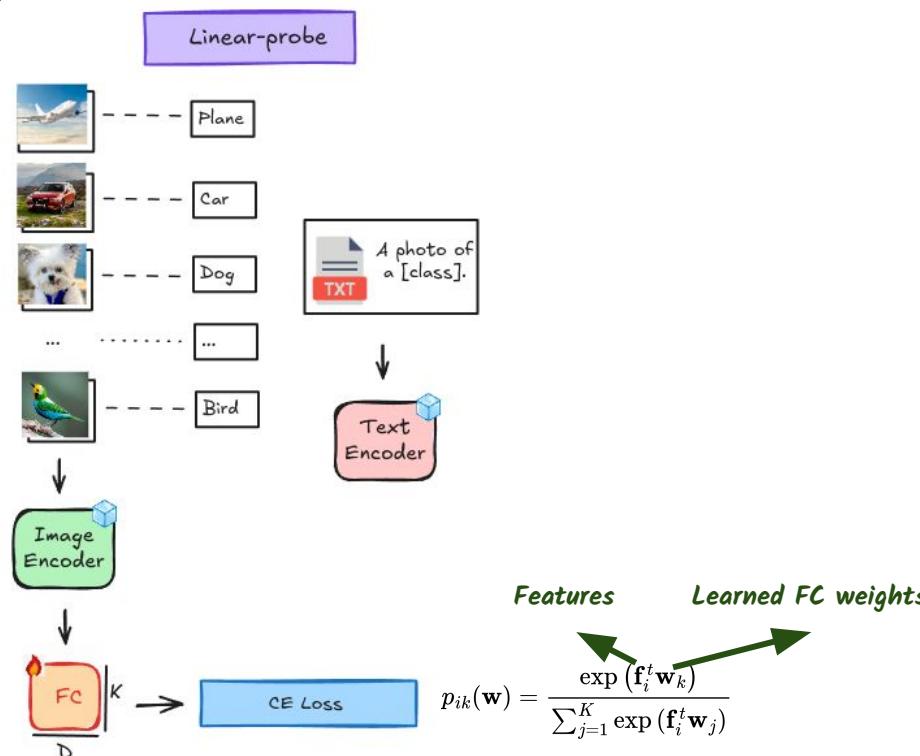
Prompt learning: Zhou, Kaiyang, et al. "Learning to prompt for vision-language models." *IJCV* (2022).

Few-Shot Adaptation (prompt learning)

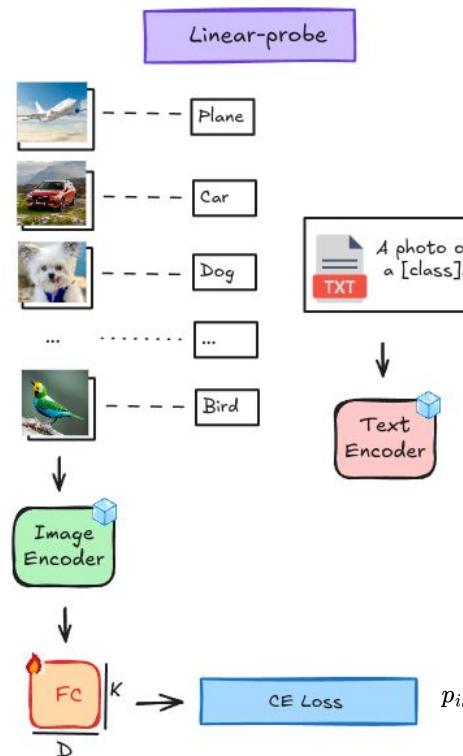


Prompt learning: Zhou, Kaiyang, et al. "Learning to prompt for vision-language models." *IJCV* (2022).
 Figure from: Huang, Y, et al. . "LP++: A Surprisingly Strong Linear Probe for Few-Shot CLIP" . *CVPR* (2024).

Few-Shot Adaptation (linear probe)



Few-Shot Adaptation (linear probe)



learnable image-text blending parameters Text Prototype

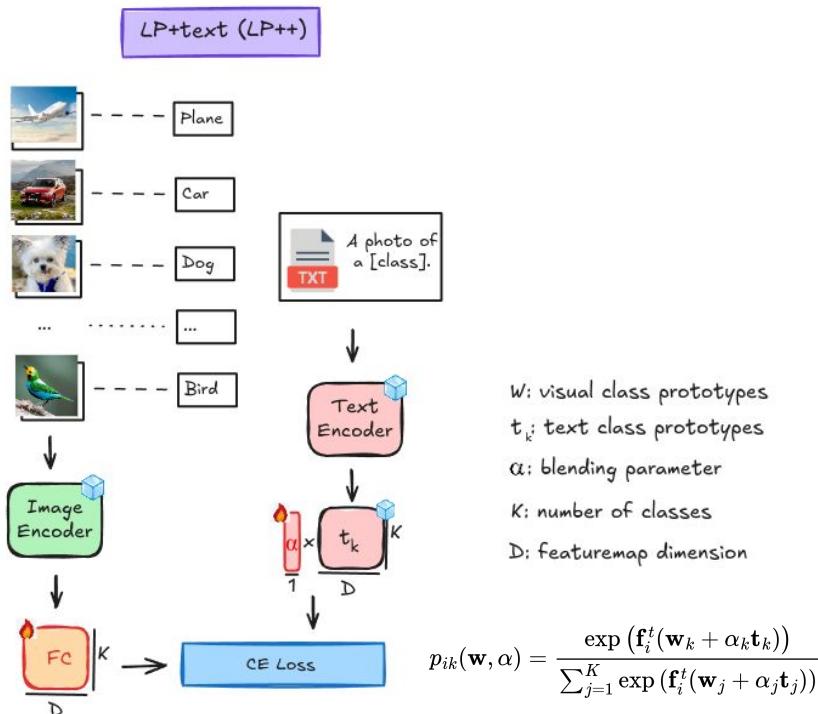
$$p_{ik}(\mathbf{w}, \alpha) = \frac{\exp (\mathbf{f}_i^t(\mathbf{w}_k + \alpha_k \mathbf{t}_k))}{\sum_{j=1}^K \exp (\mathbf{f}_i^t(\mathbf{w}_j + \alpha_j \mathbf{t}_j))}$$

Features Learned FC weights

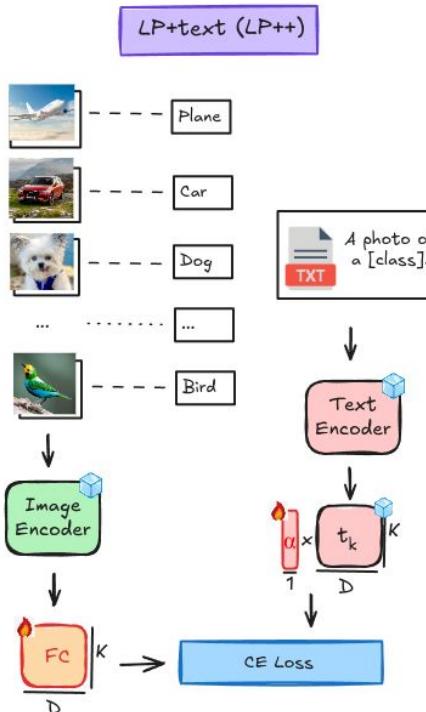
$$p_{ik}(\mathbf{w}) = \frac{\exp (\mathbf{f}_i^t \mathbf{w}_k)}{\sum_{j=1}^K \exp (\mathbf{f}_i^t \mathbf{w}_j)}$$

LP+text (LP++)

Few-Shot Adaptation (LP++)

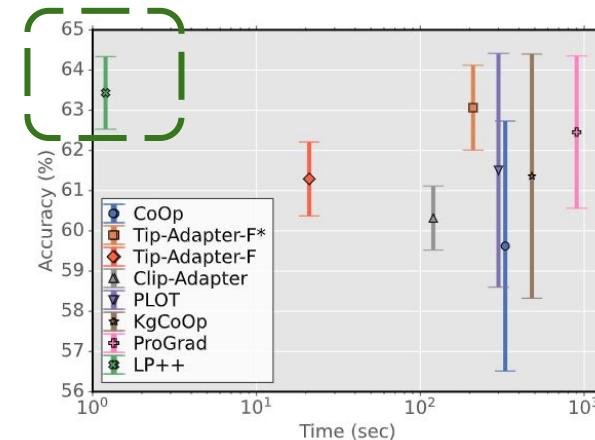


Few-Shot Adaptation (LP++)

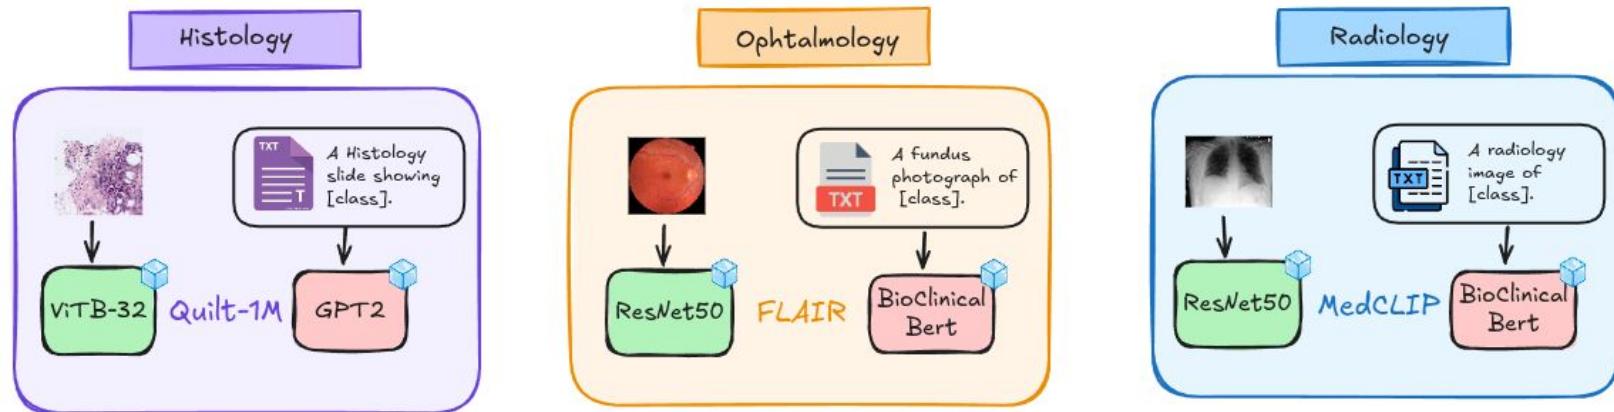


Orders-of-magnitude faster

Préserve black-box



Medical VLMs

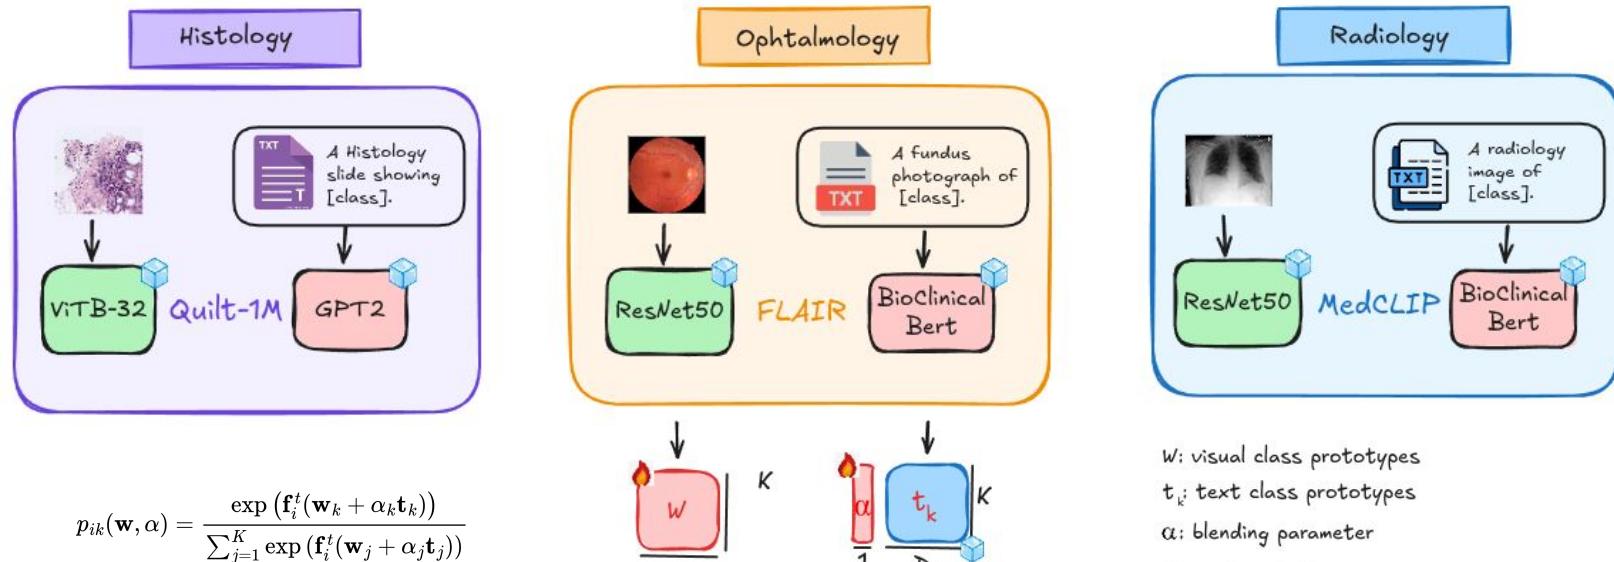


FLAIR: Silva-Rodríguez et al., A Foundation Language-Image Model of the Retina, Media 2024

Quilt-1M: Ikezogwo et al., One Million Image-Text Pairs for Histopathology, NeurIPS 2023

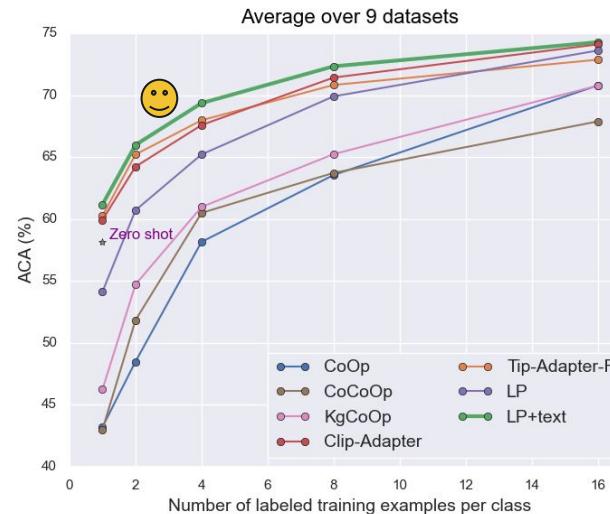
MedCLIP: Wang et al., Contrastive Learning from Unpaired medical images and text, EMNLP 2022

Few-Shot Adaptation in Medical VLMs



Few-Shot Adaptation in Medical VLMs

3 modalities / 9 datasets

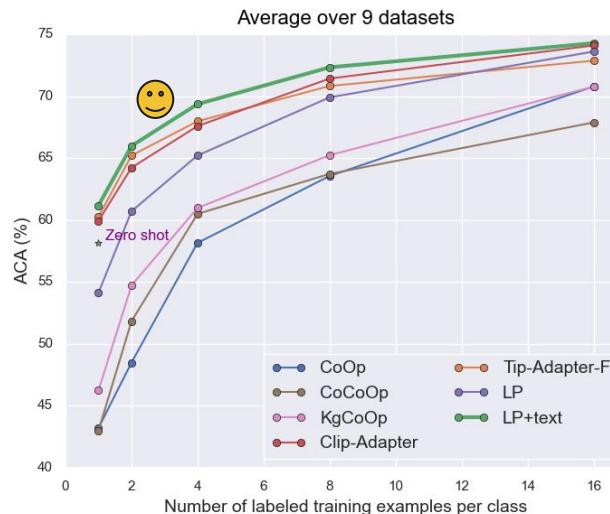


LP+text is competitive!

Shakeri, F, et al. "Few-shot adaptation of medical vision-language models." *MICCAI*, 2024.

Few-Shot Adaptation in Medical VLMs

3 modalities / 9 datasets



LP+text is competitive!

LP+text is extremely efficient!

- *Adaptation in a matter of seconds*
- *Trainable on commodity GPUs*
- *Black-box adaptation*

Methods	Category	Training Time	Black-box	#Parameters
Zero-shot [21]		n/a	✓	n/a
CoOp [35]	Prompt-Learning	3min	✗	$K \times n_{ctx1} \times D$
CoCoOp [34]		12min	✗	$n_{ctx2} \times D + C$
KgCoOp [31]		3min	✗	$K \times n_{ctx1} \times D$
Clip-Adapter [5]	CLIP-based Adapters	2min	✓	$2(D_1 \times D_2)$
Tip-adapter-F [32]		2min	✓	$K \times S \times D$
LP	Linear probe	43s	✓	$K \times D$
LP+text [7]		4s	✓	$K(D + 1)$

References

- Few-shot Adaptation of Medical Vision Language Models [Spotlight]

F Shakeri, Y Huang, JR Silva, H Bahig, A Tang, J Dolz, IB Ayed
IMedical Image Computing and Computer Assisted Intervention (MICCAI), 2024

- LP++: A Surprisingly Strong Linear Probe for Few-Shot CLIP

Y Huang, F Shakeri, J Dolz, M Boudiaf, H Bahig, IB Ayed
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024



- Boosting Vision-Language Models for Histopathology Classification: Predict all at once [Best Paper Award]

M Zanella, F Shakeri, Y Huang, H Bahig, IB Ayed
MICCAI Workshop on Foundation Models for General Medical AI (MedAGI), 2024



Questions? :)

Thank you for your attention!

